

DevSpace **Elixir Study Group**
September 2014

Working with Processes

@xavierdefrang
github.com/xavier



Erlang is designed for **massive concurrency**.

Erlang processes are **light-weight** (grow and shrink dynamically) with **small memory footprint**, **fast to create and terminate** and the **scheduling overhead is low**.

http://www.erlang.org/doc/reference_manual/processes.html

RabbitMQ Management

localhost:15672/#/

Reader

RabbitMQ™

User: guest

Cluster: rabbit@voigtkampff-2 (change)

RabbitMQ 3.3.1, Erlang R16B02

Log out

Overview

Connections

Channels

Exchanges

Queues

Admin

Visualiser

Overview

Totals

Queued messages (chart: last minute) (?)

1.0

0.0

08:40:10

08:40:20

08:40:30

Ready

Unacknowledged

0 msg

0 msg

Message rates (chart: last minute)

Currently idle

Global counts (?)

Connections: 0

Erlang processes

200

1048576 available

Memory

41MB

2.5GB high watermark

Nodes

Name	File descriptors (?)	Socket descriptors (?)	Erlang processes	Memory	Disk space	Uptime	Type
rabbit@localhost	<div>28</div> <div>256 available</div>	<div>3</div> <div>138 available</div>	<div>200</div> <div>1048576 available</div>	<div>41MB</div> <div>2.5GB high watermark</div>	<div>21GB</div> <div>48MB low watermark</div>	2m 9s	<div>Disc</div> <div>Stats</div> <div>*</div>

Ports and contexts

Listening ports

Protocol	Bound to	Port
amqp	127.0.0.1	5672
clustering	..	25672

Display a menu

Erlang Processes

- are **not** OS processes or threads
- are **cheap** (~300 memory words)
- **scheduled** by the Erlang VM
- **all Erlang / Elixir code** executes in a process

Starting a Process

```
pid = spawn(fn -> 1 + 2 end)  
#PID<0.43.0>
```

Process Identifiers

```
pid = spawn(fn -> 1 + 2 end)  
#PID<0.43.0>
```

```
Process.alive?(pid)  
false
```

```
Process.alive?(self())  
true
```

Inter-Process Communication

- shared-nothing asynchronous message passing
- each process has a **mailbox**
- simple send / receive API built-in

Example

Process 1

```
...  
send pid, {self, :sum, 1, 2}
```

```
...
```

```
receive do  
  {^pid, :ok, sum}  
  IO.puts sum  
end
```

```
...
```

Process 2

```
...
```

```
receive do  
  {sender, :sum, a, b} ->  
    send sender, {self, :ok, a+b}  
  {sender, :subtract, a, b} ->  
    send sender, {self, :ok, a-b}
```

```
...
```

```
end
```


```
...
```


Receiving Messages


- receive **blocks** until a new message matches any patterns
- unmatched messages **remain in the mailbox**

Message Order

Example: message send sequence m1, m2, m3



Message order is
guaranteed **per process**



No guarantee
that m2 will be received
after m1 and before m2

Timing Out

```
receive do
  {:hello, msg} -> msg
after
  1_000 -> "timed out after 1s"
end
```

Handling Failures

```
spawn fn -> raise "oops" end  
#PID<0.58.0>
```

Linking Processes

```
spawn_link fn -> raise "oops" end  
#PID<0.60.0>
```

```
** (EXIT from #PID<0.41.0>) an exception was raised:  
** (RuntimeError) oops  
   :erlang.apply/2
```

Linking can also be done manually using `Process.link/1`

Named Processes

```
Process.register(pid, :my_proc)
# true
send :my_proc, {:msg, self}
# ...
Process.registered
# [:elixir_sup, :error_logger, ...
  :my_proc]
Process.whereis :my_proc
# #PID<0.68.0>
```

Managing State

- Processes as key mechanism of **state management**
- Encapsulation of “persistent” state
- Mutations as response to incoming messages
- Sounds a lot like OOP, doesn't it?

State is the **value** of an **identity**
at one point in **time**

Rich Hickey

Example: Counter

```
c1 = Counter.new  
c2 = Counter.new
```

```
Counter.inc(c1)  
Counter.inc(c1)  
Counter.inc(c2)
```

```
I0.inspect Counter.value(c1)  
# 2  
I0.inspect Counter.value(c2)  
# 1
```

Counter Initialization

```
defmodule Counter do

  def new do
    spawn_link(fn -> loop(0) end)
  end

  # ...

end
```

Counter Loop

```
defp loop(counter) do
  receive do
    {:inc, _} ->
      loop(counter + 1)
    {:val, sender} ->
      send sender, {:ok, self(), counter}
      loop(counter)
  end
end
```

Incrementing

```
def inc(counter) do  
  send counter, { :inc, self }  
end
```

Querying the Value

```
def value(counter) do
  send counter, { :val, self }
  receive do
    { :ok, ^counter, val } ->
      val
  end
end
```



Caret Operator

```
iex(1)> a = 1
```

```
# 1
```

```
iex(2)> ^a = 1
```

```
# 1
```

```
iex(3)> a = 2
```

```
# 2
```

```
iex(4)> ^a = 1
```

```
** (MatchError) no match of  
    right hand side value: 1
```

The caret indicates that we want to match against the value currently bound to the variable rather than rebinding it

Nodes

- Nodes are **named Erlang runtimes**
- Code can run **transparently** on the same or a different host
- **Authentication** using a **cookie** (shared-secret)
- Node connections are:
 - **symmetric**
if A connects to B then B is also connected to A
 - **transitive** (by default)
if A connects to B and B is connected to C then A is connected to C

Starting a Node

```
$ iex --name node_name --cookie secret
```

```
$ elixir --name node_name --cookie secret
```

```
...
```


Host 1

```
[host1] $ iex --name alice --cookie secret
```

```
iex(alice@host1)1> Node.list  
[]
```

```
iex(alice@host1)2> Node.list  
[:"bob@host2"]
```

Host 2

```
[host2] $ iex --name bob --cookie secret
```

```
iex(bob@host2)1> Node.connect(:"alice@host1")  
true
```

```
iex(bob@host2)2> print_node_name =  
  fn -> IO.puts Node.self end  
#Function<20.90072148/0 in :erl_eval.expr/5>
```

```
iex(bob@host2)3>  
  Node.spawn(:"alice@host1", print_node_name)  
alice@host1  
#PID<9088.64.0>
```



Hold on a second...

```
iex(bob@host2)2> print_node_name = fn -> IO.puts Node.self end  
iex(bob@host2)3> Node.spawn(:"alice@host1", print_node_name)  
bob@host1
```

The code of `print_node_name` is executed on **host1**
but `IO.puts` writes to **host2** standard output

`IO.puts` actually communicates
with an **I/O server** running on **host2**

Code created on a given node inherits
its **process hierarchy**



Useful Abstractions

- Processes are rarely used as-is
- Larger software is designed as a tree of supervised applications using the **OTP API** (`gen_server`, `gen_fsm`, `gen_event`, `supervisor`, ...)
- The Elixir stdlib provides a handful of ready-to-use convenient abstractions for common use cases:
 - **Agent**, for state management
 - **Task**, for asynchronous workers

Further Readings

- Elixir Process module
<http://elixir-lang.org/docs/stable/elixir/Process.html>
- Elixir Node module
<http://elixir-lang.org/docs/stable/elixir/Node.html>
- How Erlang does scheduling?
<http://jlouisramblings.blogspot.be/2013/01/how-erlang-does-scheduling.html>
- Understanding the Erlang Scheduler
https://www.youtube.com/watch?v=tBAM_N9qPno

Playing Around

- In **IEx**:
 - **First:**
 - Send some messages to the **current process**
 - Receive and print those messages
 - **Then:**
 - Spawn a **new process** performing a simple calculation
 - Retrieve the result of the calculation in the IEx shell

Exercises

<https://github.com/belgian-elixir-study-group/meetup-materials>

- In meetup2:
 1. Counter
 2. Flush
 3. Parallel Map



Use cookie devspace